


Key Idea: “Logical Data Independence”


views

logical data independence



relations

physical data independence



files and pointers

```
SELECT *
FROM my_sequences
```

```
SELECT seq
FROM ncbi_sequences
WHERE seq = 'GATTACGATATTA';
```

```
f = fopen('table_file');
fseek(10030440);
while (True) {
    fread(&buf, 1, 8192, f);
    if (buf == GATTACGATATTA) {
```

What are Views?

- A **view** is just a query with a name
- We can use the view just like a real table

Why can we do this?

Because we know that every query returns a relation:
We say that the language is “algebraically closed”

View example

A view is a relation defined by a query

Purchase(customer, product, store)
Product(pname, price)

StorePrice(store, price)

```
CREATE VIEW StorePrice AS  
SELECT x.store, y.price  
FROM Purchase x, Product y  
WHERE x.pid = y.pid
```

This is like a new table
StorePrice(store,price)

Customer(cid, name, city)
Purchase(customer, product, store)
Product(pname, price)

StorePrice(store, price)

How to Use a View?

- A "high end" store is a store that sold some product over 1000. For each customer, find all the high end stores that they visit. Return a set of (customer-name, high-end-store) pairs.

```
SELECT DISTINCT z.name, u.store
FROM Customer z, Purchase u, StorePrice v
WHERE z.cid = u.customer
AND u.store = v.store
AND v.price > 1000
```

Key Idea: Indexes

- Databases are especially, but not exclusively, effective at “Needle in Haystack” problems:
 - Extracting small results from big datasets
 - Your query will **always*** finish, regardless of dataset size.
 - Indexes are easily built and automatically used when appropriate

```
CREATE INDEX seq_idx ON sequence(seq) ;
```

```
SELECT seq  
  FROM sequence  
 WHERE seq = 'GATTACGATATTA' ;
```

***almost**